

Introducing Monte Carlo Methods with R

Christian P. Robert

Université Paris Dauphine

xian@ceremade.dauphine.fr

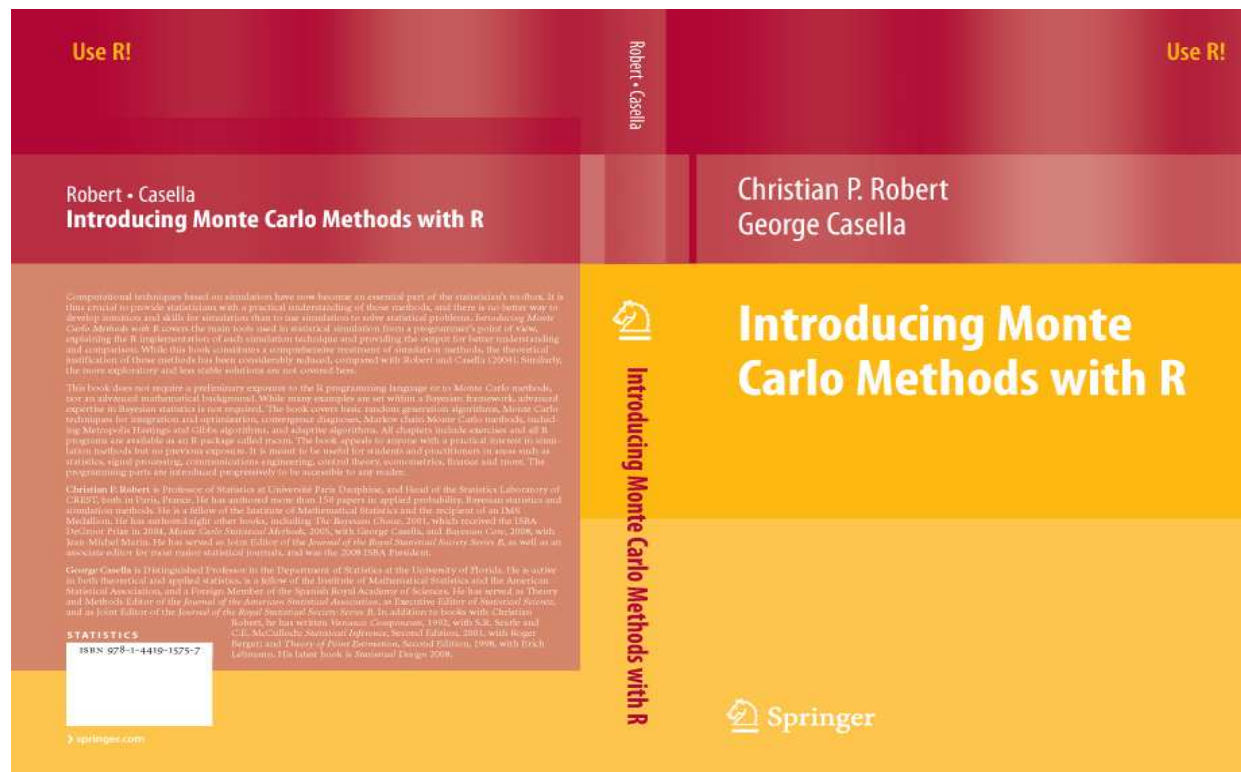
George Casella

University of Florida

casella@ufl.edu

Based on

- **Introducing Monte Carlo Methods with R**, 2009, Springer-Verlag
- Data and R programs for the course available at <http://www.stat.ufl.edu/casella/IntroMonte/>



Chapter 1: Basic R Programming

“You’re missing the big picture,” he told her. “A good album should be more than the sum of its parts.”

Ian Rankin
Exit Music

This Chapter

- ▶ We introduce the programming language R
- ▶ Input and output, data structures, and basic programming commands
- ▶ The material is both crucial and unavoidably sketchy

Basic R Programming

Introduction

- ▶ This is a quick introduction to **R**
- ▶ There are entire books devoted to **R**
 - ▷ **R Reference Card**
 - ▷ available at <http://cran.r-project.org/doc/contrib/Short-refcard.pdf>
- ▶ Take Heart!
 - ▷ The syntax of **R** is simple and logical
 - ▷ The best, and in a sense the only, way to learn **R** is through trial-and-error
- ▶ Embedded help commands **help()** and **help.search()**
 - ▷ **help.start()** opens a Web browser linked to the local manual pages

Basic R Programming

Why R ?

- ▶ There exist other languages, most (all?) of them faster than **R**, like **Matlab**, and even free, like **C** or **Python**.
- ▶ The language combines a sufficiently high power (for an interpreted language) with a very clear syntax both for statistical computation and graphics.
- ▶ **R** is a flexible language that is *object-oriented* and thus allows the manipulation of complex data structures in a condensed and efficient manner.
- ▶ Its graphical abilities are also remarkable
 - ▷ Possible interfacing with **L^AT_EX** using the package **Sweave**.

Basic R Programming

Why R ?

- ▶ R offers the additional advantages of being a free and open-source system
 - ▷ There is even an R newsletter, *R-News*
 - ▷ Numerous (free) Web-based tutorials and user's manuals
- ▶ It runs on all platforms: **Mac**, **Windows**, **Linux** and **Unix**
- ▶ R provides a powerful *interface*
 - ▷ Can integrate programs written in other languages
 - ▷ Such as **C**, **C++**, **Fortran**, **Perl**, **Python**, and **Java**.
- ▶ It is increasingly common to see people who develop new methodology simultaneously producing an R package
- ▶ Can interface with **WinBugs**

Basic R Programming

Getting started

- ▶ Type 'demo()' for some demos; `demo(image)` and `demo(graphics)`
- ▶ 'help()' for on-line help, or 'help.start()' for an HTML browser interface to help.
- ▶ Type 'q()' to quit R.
- ▶ Additional packages can be loaded via the `library` command, as in
 - > `library(combinat)` # combinatorics utilities
 - > `library(datasets)` # The R Datasets Package
 - ▷ There exist hundreds of packages available on the Web.
 - > `install.package("mcsn")`
- ▶ A `library` call is required each time R is launched

Basic R Programming

R objects

- ▶ R distinguishes between several types of *objects*
 - ▷ scalar, vector, matrix, time series, data frames, functions, or graphics.
 - ▷ An R object is mostly characterized by a *mode*
 - ▷ The different modes are
 - **null** (empty object),
 - **logical** (**TRUE** or **FALSE**),
 - **numeric** (such as **3**, **0.14159**, or **2+sqrt(3)**),
 - **complex**, (such as **3-2i** or **complex(1,4,-2)**), and
 - **character** (such as **"Blue"**, **"binomial"**, **"male"**, or **"y=a+bx"**),
- ▶ The R function **str** applied to any R object will show its structure.

Basic R Programming Interpreted

- ▶ R operates on those types as a regular function would operate on a scalar
- ▶ R is interpreted \Rightarrow Slow
- ▶ Avoid loops in favor of matrix manipulations

Basic R Programming – The **vector** class

```
> a=c(5,5.6,1,4,-5)  build the object a containing a numeric vector  
                      of dimension 5 with elements 5, 5.6, 1, 4, -5  
  
> a[1]              display the first element of a  
  
> b=a[2:4]          build the numeric vector b of dimension 3  
                      with elements 5.6, 1, 4  
  
> d=a[c(1,3,5)]     build the numeric vector d of dimension 3  
                      with elements 5, 1, -5  
  
> 2*a               multiply each element of a by 2  
                      and display the result  
  
> b%%3              provides each element of b modulo 3
```

Basic R Programming

More **vector** class

- > `e=3/d` build the numeric vector **e** of dimension 3
 and elements $3/5$, 3 , $-3/5$
- > `log(d*e)` multiply the vectors **d** and **e** term by term
 and transform each term into its natural logarithm
- > `sum(d)` calculate the sum of **d**
- > `length(d)` display the length of **d**

Basic R Programming

Even more **vector** class

<code>> t(d)</code>	transpose d , the result is a row vector
<code>> t(d)*e</code>	elementwise product between two vectors with identical lengths
<code>> t(d)%*%e</code>	matrix product between two vectors with identical lengths
<code>> g=c(sqrt(2),log(10))</code>	build the numeric vector g of dimension 2 and elements $\sqrt{2}$, $\log(10)$
<code>> e[d==5]</code>	build the subvector of e that contains the components e[i] such that d[i]=5
<code>> a[-3]</code>	create the subvector of a that contains all components of a but the third.
<code>> is.vector(d)</code>	display the logical expression TRUE if a vector and FALSE else

Basic R Programming
Comments on the **vector** class

- The ability to apply scalar functions to vectors: [Major Advantage of R](#).

- ▷ `> lgamma(c(3,5,7))`

- ▷ returns the vector with components $(\log \Gamma(3), \log \Gamma(5), \log \Gamma(7))$.

- Functions that are specially designed for vectors include

sample, **permn**, **order**, **sort**, and **rank**

- ▷ All manipulate the order in which the components of the vector occur.

- ▷ **permn** is part of the **combinat** library

- The components of a vector can also be identified by names.

- ▷ For a vector **x**, **names(x)** is a vector of characters of the same length as **x**

Basic R Programming

The `matrix`, `array`, and `factor` classes

- ▶ The `matrix` class provides the R representation of matrices.
- ▶ A typical entry is
 - > `x=matrix(vec,nrow=n,ncol=p)`
 - ▷ Creates an $n \times p$ matrix whose elements are of the dimension np vector `vec`
- ▶ Some manipulations on matrices
 - ▷ The standard matrix product is denoted by `%*%`,
 - ▷ while `*` represents the term-by-term product.
 - ▷ `diag` gives the vector of the diagonal elements of a matrix
 - ▷ `crossprod` replaces the product `t(x)%*%y` on either vectors or matrices
 - ▷ `crossprod(x,y)` more efficient
 - ▷ `apply` is easy to use for functions operating on matrices by row or column

Basic R Programming

Some **matrix** commands

<code>> x1=matrix(1:20,nrow=5)</code>	build the numeric matrix x1 of dimension 5×4 with first row 1, 6, 11, 16
<code>> x2=matrix(1:20,nrow=5,byrow=T)</code>	build the numeric matrix x2 of dimension 5×4 with first row 1, 2, 3, 4
<code>> a=x1%*%t(x2)</code>	matrix product
<code>> c=x1*x2</code>	term-by-term product between x1 and x2
<code>> dim(x1)</code>	display the dimensions of x1
<code>> b[,2]</code>	select the second column of b
<code>> b[c(3,4),]</code>	select the third and fourth rows of b
<code>> b[-2,]</code>	delete the second row of b
<code>> rbind(x1,x2)</code>	vertical merging of x1 and x2 <code>rbind(*)rbind</code>
<code>> cbind(x1,x2)</code>	horizontal merging of x1 and x2 <code>cbind(*)cbind</code>
<code>> apply(x1,1,sum)</code>	calculate the sum of each row of x1
<code>> as.matrix(1:10)</code>	turn the vector 1:10 into a 10×1 matrix

► Lots of other commands that we will see throughout the course

Basic R Programming
The `list` and `data.frame` classes
The Last One

- ▶ A **list** is a collection of arbitrary objects known as its *components*
 - > `li=list(num=1:5,y="color",a=T)` create a list with three arguments
 - ▶ The last class we briefly mention is the **data frame**
 - ▷ A list whose elements are possibly made of differing modes and attributes
 - ▷ But have the same length
- | | |
|---|---|
| <ul style="list-style-type: none">> <code>v1=sample(1:12,30,rep=T)</code>> <code>v2=sample(LETTERS[1:10],30,rep=T)</code>> <code>v3=runif(30)</code>> <code>v4=rnorm(30)</code>> <code>xx=data.frame(v1,v2,v3,v4)</code> | <ul style="list-style-type: none">simulate 30 independent uniform $\{1, 2, \dots, 12\}$simulate 30 independent uniform $\{a, b, \dots, j\}$simulate 30 independent uniform $[0, 1]$simulate 30 independent standard normalscreate a data frame |
|---|---|
- ▶ R **code**

Probability distributions in R

- R , or the web, has about all probability distributions
- Prefixes: **p**, **d**, **q**, **r**

Distribution	Core	Parameters	Default Values
Beta	beta	shape1, shape2	
Binomial	binom	size, prob	
Cauchy	cauchy	location, scale	0, 1
Chi-square	chisq	df	
Exponential	exp	1/mean	1
F	f	df1, df2	
Gamma	gamma	shape, 1/scale	NA, 1
Geometric	geom	prob	
Hypergeometric	hyper	m, n, k	
Log-normal	lnorm	mean, sd	0, 1
Logistic	logis	location, scale	0, 1
Normal	norm	mean, sd	0, 1
Poisson	pois	lambda	
Student	t	df	
Uniform	unif	min, max	0, 1
Weibull	weibull	shape	

Basic and not-so-basic statistics

t -test

► Testing equality of two means

```
> x=rnorm(25) #produces a N(0,1) sample of size 25
> t.test(x)
```

One Sample t -test

```
data:  x
t = -0.8168, df = 24, p-value = 0.4220
alternative hypothesis: true mean is not equal to 0
95 percent confidence interval:
 -0.4915103  0.2127705
sample estimates:
mean of x
-0.1393699
```

Basic and not-so-basic statistics

Correlation

► Correlation

```
> attach(faithful) #resident dataset  
> cor.test(faithful[,1],faithful[,2])
```

Pearson's product-moment correlation

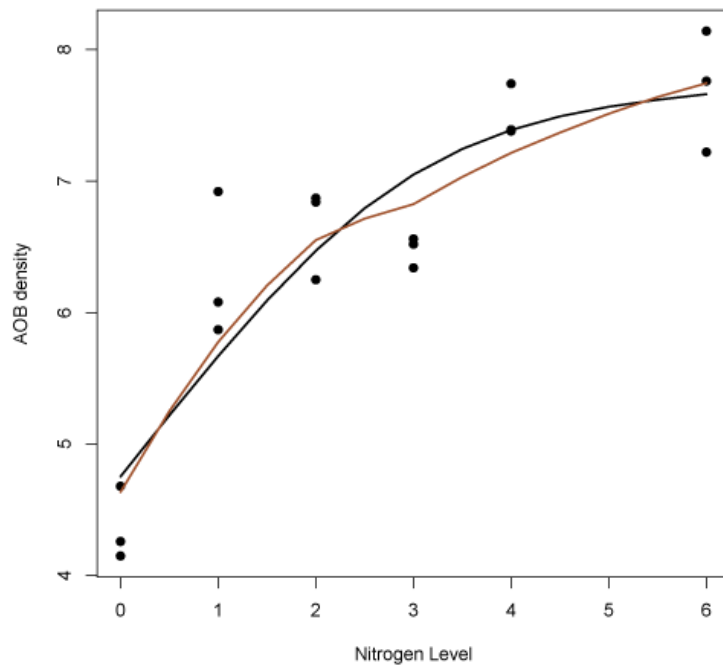
```
data: faithful[, 1] and faithful[, 2]  
t = 34.089, df = 270, p-value < 2.2e-16  
alternative hypothesis: true correlation is not equal to 0  
95 percent confidence interval:  
 0.8756964 0.9210652  
sample estimates:  
      cor  
0.9008112
```

► R [code](#)

Basic and not-so-basic statistics

Splines

- ▶ Nonparametric regression with **loess** function or using *natural splines*
- ▶ Relationship between nitrogen level in soil and abundance of a bacteria **AOB**



- ▶ Natural spline fit (*dark*)
 - ▷ With **ns=2** (linear model)
- ▶ Loess fit (*brown*) with **span=1.25**
- ▶ R **code**

Basic and not-so-basic statistics Generalized Linear Models

- Fitting a binomial (logistic) glm to the probability of suffering from diabetes for a woman within the Pima Indian population

```
> glm(formula = type ~ bmi + age, family = "binomial", data = Pima.tr)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.7935	-0.8368	-0.5033	1.0211	2.2531

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-6.49870	1.17459	-5.533	3.15e-08 ***
bmi	0.10519	0.02956	3.558	0.000373 ***
age	0.07104	0.01538	4.620	3.84e-06 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 256.41 on 199 degrees of freedom

Residual deviance: 215.93 on 197 degrees of freedom

AIC: 221.93

Number of Fisher Scoring iterations: 4

Basic and not-so-basic statistics
Generalized Linear Models – Comments

- ▶ Concluding with the significance both of the body mass index **bmi** and the age
- ▶ Other generalized linear models can be defined by using a different **family** value
 - > `glm(y ~x, family=quasi(var="mu^2", link="log"))`
 - ▷ Quasi-Likelihood also
- ▶ Many many other procedures
 - ▷ Time series, anova,...
- ▶ One last one

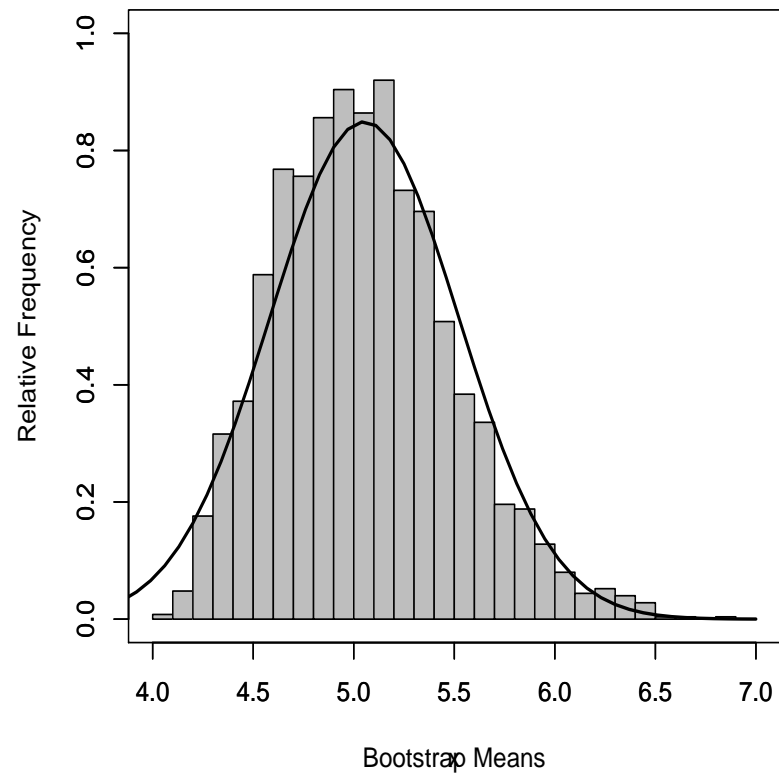
Basic and not-so-basic statistics

Bootstrap

- ▶ The bootstrap procedure uses the empirical distribution as a substitute for the true distribution to construct variance estimates and confidence intervals.
 - ▷ A sample X_1, \dots, X_n is **resampled** with replacement
 - ▷ The empirical distribution has a finite but large support made of n^n points
- ▶ For example, with data y , we can create a bootstrap sample y^* using the code
 - > `ystar=sample(y,replace=T)`
 - ▷ For each resample, we can calculate a mean, variance, etc

Basic and not-so-basic statistics

Simple illustration of bootstrap



- ▶ A histogram of 2500 bootstrap means
- ▶ Along with the normal approximation
- ▶ Bootstrap shows some skewness
- ▶ R `code`

Basic and not-so-basic statistics
Bootstrapping Regression

- ▶ The bootstrap is not a panacea
 - ▷ Not always clear which quantity should be bootstrapped
 - ▷ In regression, bootstrapping the residuals is preferred

- ▶ Linear regression

$$Y_{ij} = \alpha + \beta x_i + \varepsilon_{ij},$$

α and β are the unknown intercept and slope, ε_{ij} are the iid normal errors

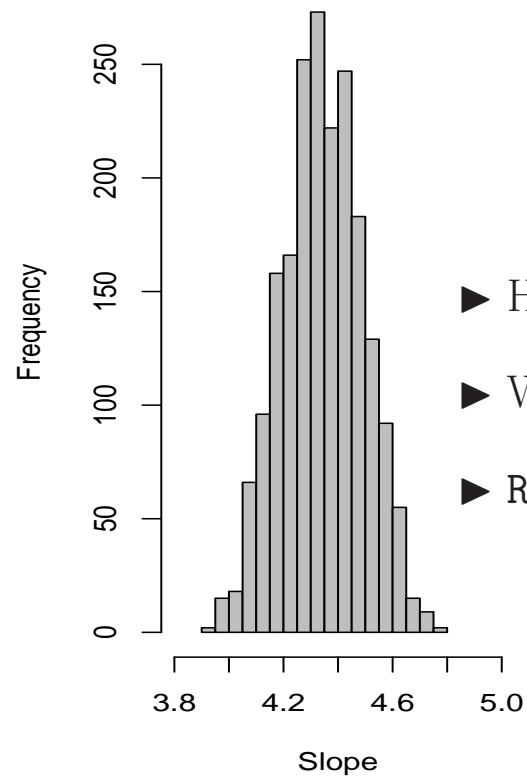
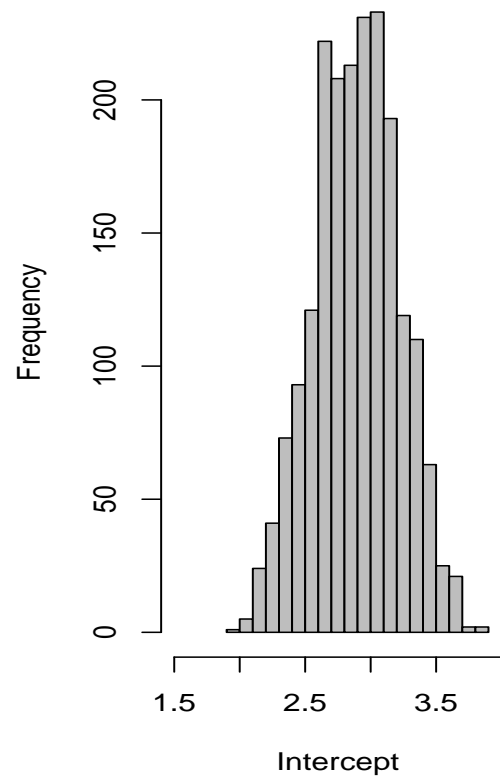
- ▶ The residuals from the least squares fit are given by

$$\hat{\varepsilon}_{ij} = y_{ij} - \hat{\alpha} - \hat{\beta}x_i,$$

- ▷ We bootstrap the residuals
- ▷ Produce a new sample $(\hat{\varepsilon}_{ij}^*)_{ij}$ by resampling from the $\hat{\varepsilon}_{ij}$'s
- ▷ The bootstrap samples are then $y_{ij}^* = y_{ij} + \hat{\varepsilon}_{ij}^*$

Basic and not-so-basic statistics

Bootstrapping Regression – 2



- Histogram of 2000 bootstrap samples
- We can also get confidence intervals
- R `code`

Basic R Programming Some Other Stuff

- ▶ Graphical facilities
 - ▷ Can do a lot; see `plot` and `par`
- ▶ Writing new R functions
 - ▷ `h=function(x)(sin(x)^2+cos(x)^3)^(3/2)`
 - ▷ We will do this a lot
- ▶ Input and output in R
 - ▷ `write.table`, `read.table`, `scan`
- ▶ Don't forget the `mcs` package

Chapter 2: Random Variable Generation

“It has long been an axiom of mine that the little things are infinitely the most important.”

Arthur Conan Doyle
A Case of Identity

This Chapter

- ▶ We present practical techniques that can produce random variables
- ▶ From both standard and nonstandard distributions
- ▶ First: Transformation methods
- ▶ Next: Indirect Methods - Accept–Reject

Introduction

- ▶ Monte Carlo methods rely on
 - ▷ The possibility of producing a supposedly endless flow of random variables
 - ▷ For well-known or new distributions.
- ▶ Such a simulation is, in turn,
 - ▷ Based on the production of uniform random variables on the interval $(0, 1)$.
- ▶ We are not concerned with the [details](#) of producing uniform random variables
- ▶ We assume the existence of such a sequence

Introduction

Using the R Generators

R has a large number of functions that will generate the standard random variables

```
> rgamma(3,2.5,4.5)
```

produces three independent generations from a $\mathcal{G}(5/2, 9/2)$ distribution

- ▶ It is therefore,
 - ▷ Counter-productive
 - ▷ Inefficient
 - ▷ And even dangerous,
- ▶ To generate from those standard distributions
- ▶ If it is built into R , use it
- ▶ But....we will practice on these.
- ▶ The principles are essential to deal with distributions that are not built into R.

Uniform Simulation

- ▶ The uniform generator in R is the function `runif`
- ▶ The only required entry is the number of values to be generated.
- ▶ The other optional parameters are `min` and `max`, with R **code**

```
> runif(100, min=2, max=5)
```

will produce 100 random variables $\mathcal{U}(2, 5)$.

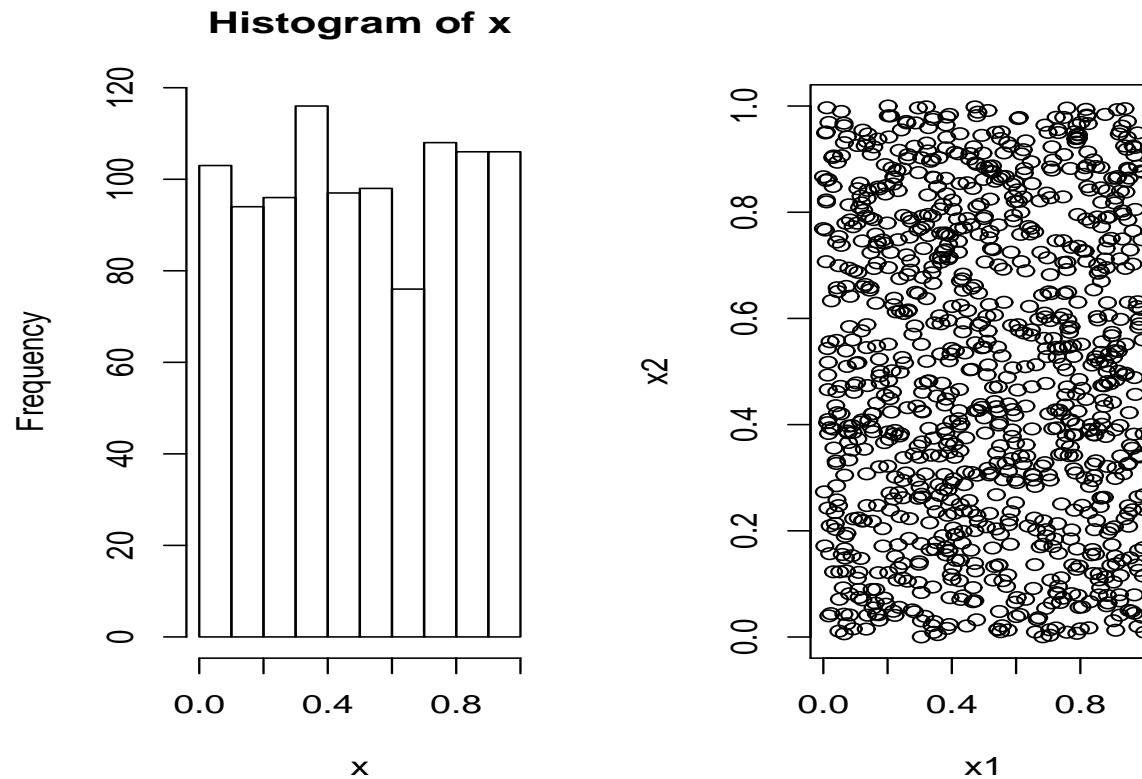
Uniform Simulation Checking the Generator

- ▶ A quick check on the properties of this uniform generator is to
 - ▷ Look at a histogram of the X_i 's,
 - ▷ Plot the pairs (X_i, X_{i+1})
 - ▷ Look at the estimate autocorrelation function
- ▶ Look at the **R code**

```

> Nsim=10^4           #number of random numbers
> x=runif(Nsim)
> x1=x[-Nsim]         #vectors to plot
> x2=x[-1]            #adjacent pairs
> par(mfrow=c(1,3))
> hist(x)
> plot(x1,x2)
> acf(x)
    
```


Uniform Simulation
Plots from the Generator



- Histogram (*left*), pairwise plot (*center*), and estimated autocorrelation function (*right*) of a sequence of 10^4 uniform random numbers generated by `runif`.

Uniform Simulation

Some Comments

- ▶ Remember: `runif` does not involve randomness per se.
- ▶ It is a deterministic sequence based on a random starting point.
- ▶ The R function `set.seed` can produce the same sequence.

```
> set.seed(1)
> runif(5)
[1] 0.2655087 0.3721239 0.5728534 0.9082078 0.2016819
> set.seed(1)
> runif(5)
[1] 0.2655087 0.3721239 0.5728534 0.9082078 0.2016819
> set.seed(2)
> runif(5)
[1] 0.0693609 0.8177752 0.9426217 0.2693818 0.1693481
```

- ▶ Setting the seed determines all the subsequent values

The Inverse Transform

► The *Probability Integral Transform*

▷ Allows us to transform a uniform into any random variable

► For example, if X has density f and cdf F , then we have the relation

$$F(x) = \int_{-\infty}^x f(t) dt,$$

and we set $U = F(X)$ and solve for X


► Example 2.1

▷ If $X \sim \mathcal{Exp}(1)$, then $F(x) = 1 - e^{-x}$

▷ Solving for x in $u = 1 - e^{-x}$ gives $x = -\log(1 - u)$

Generating Exponentials

```
> Nsim=10^4           #number of random variables
> U=runif(Nsim)
> X=-log(U)           #transforms of uniforms
> Y=rexp(Nsim)         #exponentials from R
> par(mfrow=c(1,2))   #plots
> hist(X,freq=F,main="Exp from Uniform")
> hist(Y,freq=F,main="Exp from R")
```



- Histograms of exponential random variables
 - ▷ Inverse transform (*right*)
 - ▷ R command **rexp** (*left*)
 - ▷ $\mathcal{Exp}(1)$ density on top

Generating Other Random Variables From Uniforms

- ▶ This method is useful for other probability distributions
 - ▷ Ones obtained as a transformation of uniform random variables
- ▶ Logistic pdf: $f(x) = \frac{1}{\beta} \frac{e^{-(x-\mu)/\beta}}{[1+e^{-(x-\mu)/\beta}]^2}$, cdf: $F(x) = \frac{1}{1+e^{-(x-\mu)/\beta}}$.
- ▶ Cauchy pdf: $f(x) = \frac{1}{\pi\sigma} \frac{1}{1+\left(\frac{x-\mu}{\sigma}\right)^2}$, cdf: $F(x) = \frac{1}{2} + \frac{1}{\pi} \arctan((x - \mu)/\sigma)$.

General Transformation Methods

- ▶ When a density f is linked in a relatively simple way
 - ▷ To another distribution easy to simulate
 - ▷ This relationship can be use to construct an algorithm to simulate from f
- ▶ If the X_i 's are iid $\mathcal{Exp}(1)$ random variables,
 - ▷ Three standard distributions can be derived as

$$Y = 2 \sum_{j=1}^{\nu} X_j \sim \chi_{2\nu}^2, \quad \nu \in \mathbb{N}^*,$$

$$Y = \beta \sum_{j=1}^a X_j \sim \mathcal{G}(a, \beta), \quad a \in \mathbb{N}^*,$$

$$Y = \frac{\sum_{j=1}^a X_j}{\sum_{j=1}^{a+b} X_j} \sim \mathcal{Be}(a, b), \quad a, b \in \mathbb{N}^*,$$

where $\mathbb{N}^* = \{1, 2, \dots\}$.

General Transformation Methods

χ^2_6 Random Variables

- For example, to generate χ^2_6 random variables, we could use the R [code](#)

```
> U=runif(3*10^4)
> U=matrix(data=U,nrow=3) #matrix for sums
> X=-log(U)                #uniform to exponential
> X=2* apply(X,2,sum)       #sum up to get chi squares
```

- [Not nearly as efficient as calling rchisq](#), as can be checked by the R [code](#)

```
> system.time(test1());system.time(test2())
  user  system elapsed
0.104   0.000   0.107
  user  system elapsed
0.004   0.000   0.004
```

- `test1` corresponds to the R [code](#) above
- `test2` corresponds to `X=rchisq(10^4,df=6)`

General Transformation Methods

Comments

- ▶ These transformations are quite simple and will be used in our illustrations.
- ▶ However, there are limits to their usefulness,
 - ▷ No odd degrees of freedom
 - ▷ No normals
- ▶ For any specific distribution, efficient algorithms have been developed.
- ▶ Thus, if **R** has a distribution built in, it is almost always worth using

General Transformation Methods A Normal Generator

- ▶ Box–Muller algorithm - two normals from two uniforms

- ▶ If U_1 and U_2 are iid $\mathcal{U}_{[0,1]}$

- ▶ The variables X_1 and X_2

$$X_1 = \sqrt{-2 \log(U_1)} \cos(2\pi U_2) , \quad X_2 = \sqrt{-2 \log(U_1)} \sin(2\pi U_2) ,$$

- ▶ Are iid $\mathcal{N}(0, 1)$ by virtue of a change of variable argument.

- ▶ The Box–Muller algorithm is exact, not a crude CLT-based approximation

- ▶ Note that this is *not* the generator implemented in **R**

- ▷ It uses the probability inverse transform

- ▷ With a very accurate representation of the normal cdf

General Transformation Methods Multivariate Normals

- ▶ Can simulate a multivariate normal variable using univariate normals
 - ▷ Cholesky decomposition of $\Sigma = AA'$
 - ▷ $Y \sim \mathcal{N}_p(0, I) \Rightarrow AY \sim \mathcal{N}_p(0, \Sigma)$
 - ▶ There is an **R** package that replicates those steps, called **rmnorm**
 - ▷ In the **mnormt** library
 - ▷ Can also calculate the probability of hypercubes with the function **sadmvn**
- ```
> sadmvn(low=c(1,2,3),upp=c(10,11,12),mean=rep(0,3),var=B)
[1] 9.012408e-05
attr(,"error")
[1] 1.729111e-08
```
- ▶ B is a positive-definite matrix
  - ▶ This is quite useful since the analytic derivation of this probability is almost always impossible.

## Discrete Distributions

- ▶ To generate discrete random variables we have an “all-purpose” algorithm.
- ▶ Based on the inverse transform principle
- ▶ To generate  $X \sim P_\theta$ , where  $P_\theta$  is supported by the integers,
  - ▷ We can calculate—the probabilities
  - ▷ Once for all, assuming we can store them

$$p_0 = P_\theta(X \leq 0), \quad p_1 = P_\theta(X \leq 1), \quad p_2 = P_\theta(X \leq 2), \quad \dots,$$

- ▷ And then generate  $U \sim \mathcal{U}_{[0,1]}$  and take

$$X = k \text{ if } p_{k-1} < U < p_k.$$

## Discrete Distributions

### Binomial

► **Example** To generate  $X \sim \mathcal{Bin}(10, .3)$

▷ The probability values are obtained by `pbinom(k, 10, .3)`

$$p_0 = 0.028, \quad p_1 = 0.149, \quad p_2 = 0.382, \dots, p_{10} = 1,$$

▷ And to generate  $X \sim \mathcal{P}(7)$ , take

$$p_0 = 0.0009, \quad p_1 = 0.0073, \quad p_2 = 0.0296, \dots,$$

▷ Stopping the sequence when it reaches 1 with a given number of decimals.

▷ For instance,  $p_{20} = 0.999985$ .

► Check the **R code**

## Discrete Distributions Comments

- ▶ Specific algorithms are usually more efficient
- ▶ Improvement can come from a judicious choice of the probabilities first computed.
- ▶ For example, if we want to generate from a Poisson with  $\lambda = 100$ 
  - ▷ The algorithm above is woefully inefficient
  - ▷ We expect most of our observations to be in the interval  $\lambda \pm 3\sqrt{\lambda}$
  - ▷ For  $\lambda = 100$  this interval is  $(70, 130)$
  - ▷ Thus, starting at 0 is quite wasteful
- ▶ A first remedy is to “ignore” what is outside of a highly likely interval
  - ▷ In the current example  $P(X < 70) + P(X > 130) = 0.00268$ .

## Discrete Distributions

### Poisson R Code

► R code that can be used to generate Poisson random variables for large values of lambda.

► The sequence `t` contains the integer values in the range around the mean.

```
> Nsim=10^4; lambda=100
> spread=3*sqrt(lambda)
> t=round(seq(max(0,lambda-spread),lambda+spread,1))
> prob=ppois(t, lambda)
> X=rep(0,Nsim)
> for (i in 1:Nsim){
+ u=runif(1)
+ X[i]=t[1]+sum(prob<u)-1 }
```

► The last line of the program checks to see what interval the uniform random variable fell in and assigns the correct Poisson value to  $X$ .

## Discrete Distributions

### Comments

- ▶ Another remedy is to start the cumulative probabilities at the mode of the discrete distribution
  - ▶ Then explore neighboring values until the cumulative probability is almost 1.
- 
- ▶ Specific algorithms exist for almost any distribution and are often quite fast.
  - ▶ So, if **R** has it, use it.
  - ▶ But **R** does not handle every distribution that we will need,

## Mixture Representations

- It is sometimes the case that a probability distribution can be naturally represented as a *mixture distribution*
- That is, we can write it in the form

$$f(x) = \int_{\mathcal{Y}} g(x|y)p(y) \, dy \quad \text{or} \quad f(x) = \sum_{i \in \mathcal{Y}} p_i f_i(x) ,$$

- ▷ The mixing distribution can be continuous or discrete.
- To generate a random variable  $X$  using such a representation,
  - ▷ we can first generate a variable  $Y$  from the mixing distribution
  - ▷ Then generate  $X$  from the selected conditional distribution



## Mixture Representations

### Generating the Mixture

#### ► Continuous

$$f(x) = \int_{\mathcal{Y}} g(x|y)p(y) \, dy \Rightarrow y \sim p(y) \text{ and } X \sim f(x|y), \text{ then } X \sim f(x)$$

#### ► Discrete

$$f(x) = \sum_{i \in \mathcal{Y}} p_i f_i(x) \Rightarrow i \sim p_i \text{ and } X \sim f_i(x), \text{ then } X \sim f(x)$$

#### ► Discrete Normal Mixture R **code**

$$\triangleright p_1 * N(\mu_1, \sigma_1) + p_2 * N(\mu_2, \sigma_2) + p_3 * N(\mu_3, \sigma_3)$$

## Mixture Representations

### Continuous Mixtures

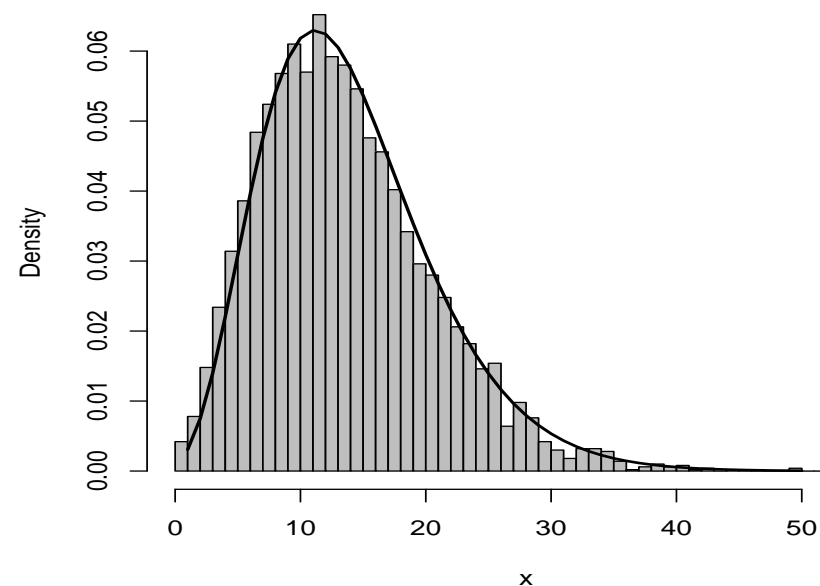
► Student's  $t$  density with  $\nu$  degrees of freedom

$$X|y \sim \mathcal{N}(0, \nu/y) \quad \text{and} \quad Y \sim \chi_\nu^2.$$

- ▷ Generate from a  $\chi_\nu^2$  then from the corresponding normal distribution
- ▷ Obviously, using **rt** is slightly more efficient

► If  $X$  is negative binomial  $X \sim \mathcal{Neg}(n, p)$

- ▷  $X|y \sim \mathcal{P}(y)$  and  $Y \sim \mathcal{G}(n, \beta)$ ,
- ▷ R **code** generates from this mixture



## Accept–Reject Methods

### Introduction

- ▶ There are many distributions where transform methods fail
- ▶ For these cases, we must turn to *indirect* methods
  - ▷ We generate a candidate random variable
  - ▷ Only accept it subject to passing a test
- ▶ This class of methods is extremely powerful.
  - ▷ It will allow us to simulate from virtually any distribution.
- ▶ Accept–Reject Methods
  - ▷ Only require the functional form of the density  $f$  of interest
  - ▷  $f$  = target,  $g$ =candidate
- ▶ Where it is simpler to simulate random variables from  $g$

## Accept–Reject Methods Accept–Reject Algorithm

- ▶ The only constraints we impose on this candidate density  $g$ 
  - ▷  $f$  and  $g$  have compatible supports (i.e.,  $g(x) > 0$  when  $f(x) > 0$ ).
  - ▷ There is a constant  $M$  with  $f(x)/g(x) \leq M$  for all  $x$ .
- ▶  $X \sim f$  can be simulated as follows.
  - ▷ Generate  $Y \sim g$  and, independently, generate  $U \sim \mathcal{U}_{[0,1]}$ .
  - ▷ If  $U \leq \frac{1}{M} \frac{f(Y)}{g(Y)}$ , set  $X = Y$ .
  - ▷ If the inequality is not satisfied, we then discard  $Y$  and  $U$  and start again.
- ▶ Note that  $M = \sup_x \frac{f(x)}{g(x)}$
- ▶  $P(\text{Accept}) = \frac{1}{M}$ , Expected Waiting Time =  $M$

## Accept–Reject Algorithm

### R Implementation

Succinctly, the Accept–Reject Algorithm is

#### Accept–Reject Method

1. Generate  $Y \sim g$ ,  $U \sim \mathcal{U}_{[0,1]}$ ;
2. Accept  $X = Y$  if  $U \leq f(Y)/Mg(Y)$ ;
3. Return to 1 otherwise.

► R implementation: If **randg** generates from  $g$

```
> u=runif(1)*M
> y=randg(1)
> while (u>f(y)/g(y))
{
 u=runif(1)*M
 y=randg(1)
}
```

► Produces a single generation  $y$  from  $f$

## Accept–Reject Algorithm Normals from Double Exponentials

► **Candidate**  $Y \sim \frac{1}{2} \exp(-|y|)$

► **Target**  $X \sim \frac{1}{\sqrt{2\pi}} \exp(-x^2/2)$

$$\frac{\frac{1}{\sqrt{2\pi}} \exp(-y^2/2)}{\frac{1}{2} \exp(-|y|)} \leq \frac{2}{\sqrt{2\pi}} \exp(1/2)$$

▷ Maximum at  $y = 1$

► Accept  $Y$  if  $U \leq \exp(-.5Y^2 + |Y| - .5)$

► Look at R **code**

## Accept–Reject Algorithm Theory

- ▶ Why does this method work?
- ▶ A straightforward probability calculation shows

$$P(Y \leq x | \text{Accept}) = P(Y \leq x | U \leq f(Y)/\{Mg(Y)\}) = P(X \leq x)$$

▷ Simulating from  $g$ , the output of this algorithm is exactly distributed from  $f$ .





- ▶ The Accept–Reject method is applicable in any dimension
- ▶ As long as  $g$  is a density over the same space as  $f$ .

- ▶ Only need to know  $f/g$  up to a constant
- ▶ Only need an upper bound on  $M$

## Accept–Reject Algorithm Betas from Uniforms

- Generate  $X \sim \text{beta}(a, b)$ .
- No direct method if  $a$  and  $b$  are not integers.
- Use a uniform candidate
- For  $a = 2.7$  and  $b = 6.3$

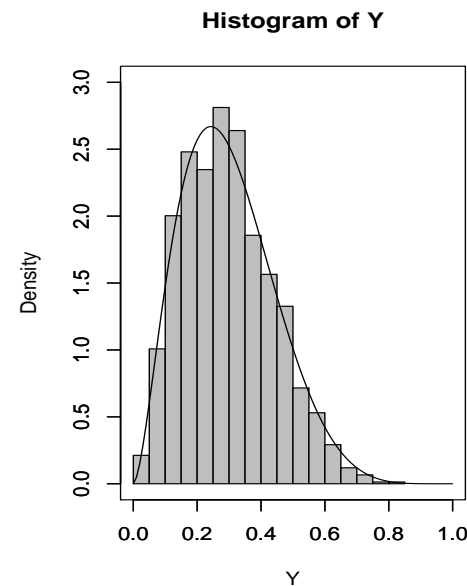
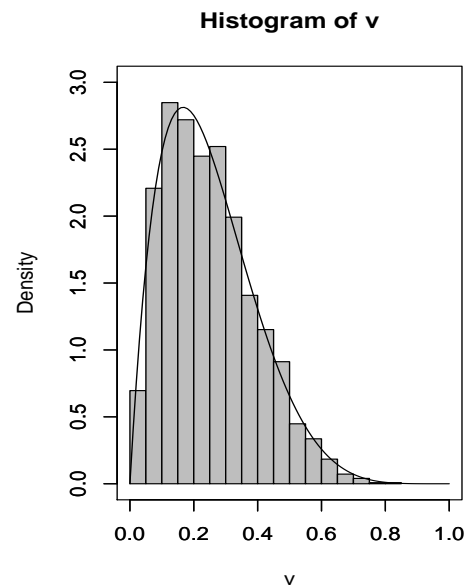


► Acceptance Rate = 37%



## Accept–Reject Algorithm Betas from Betas

- Generate  $X \sim \text{beta}(a, b)$ .
- No direct method if  $a$  and  $b$  are not integers.
- Use a beta candidate
- For  $a = 2.7$  and  $b = 6.3$ ,  $Y \sim \text{beta}(2, 6)$



► Acceptance Rate = 60%

Accept–Reject Algorithm  
Betas from Betas-Details

- ▶ Beta density  $\propto x^a(1-x)^b$
- ▶ Can generate if  $a$  and  $b$  integers
- ▶ If not, use candidate with  $a_1$  and  $b_1$  integers

$$\frac{y^a(1-y)^b}{y^{a_1}(1-y)^{b_1}} \text{ maximized at } y = \frac{a - a_1}{a - a_1 + b - b_1}$$

- ▷ Need  $a_1 < a$  and  $b_1 < b$
- ▶ Efficiency  $\uparrow$  as the candidate gets closer to the target
- ▶ Look at R code

## Accept–Reject Algorithm

### Comments

⚡ Some key properties of the Accept–Reject algorithm::

1. Only the ratio  $f/M$  is needed

▷ So the algorithm does not depend on the normalizing constant.

2. The bound  $f \leq Mg$  need not be tight

▷ Accept–Reject is valid, but less efficient, if  $M$  is replaced with a larger constant.

3. The probability of acceptance is  $1/M$

▷ So  $M$  should be as small as possible for a given computational effort.

## Chapter 3: Monte Carlo Integration

*“Every time I think I know what’s going on, suddenly there’s another layer of complications. I just want this damn thing solved.”*

**John Scalzi**  
*The Last Colony*

### **This Chapter**

- ▶ This chapter introduces the major concepts of Monte Carlo methods
- ▶ The validity of Monte Carlo approximations relies on the Law of Large Numbers
- ▶ The versatility of the representation of an integral as an expectation

## Monte Carlo Integration

### Introduction

- We will be concerned with evaluating integrals of the form

$$\int_{\mathcal{X}} h(x) f(x) \, dx,$$

- ▷  $f$  is a density
  - ▷ We can produce an almost infinite number of random variables from  $f$
- We apply probabilistic results
  - ▷ Law of Large Numbers
  - ▷ Central Limit Theorem
- The Alternative - Deterministic Numerical Integration
  - ▷ R functions `area` and `integrate`
  - ▷ OK in low (one) dimensions
  - ▷ Usually needs some knowledge of the function

## Classical Monte Carlo Integration

### The Monte Carlo Method

- The generic problem: Evaluate

$$\mathbb{E}_f[h(X)] = \int_{\mathcal{X}} h(x) f(x) \, dx,$$

- ▷  $X$  takes its values in  $\mathcal{X}$

- The Monte Carlo Method

- ▷ Generate a sample  $(X_1, \dots, X_n)$  from the density  $f$
  - ▷ Approximate the integral with

$$\bar{h}_n = \frac{1}{n} \sum_{j=1}^n h(x_j) ,$$

## Classical Monte Carlo Integration

### Validating the Monte Carlo Method

#### ► The Convergence

$$\bar{h}_n = \frac{1}{n} \sum_{j=1}^n h(x_j) \rightarrow \int_{\mathcal{X}} h(x) f(x) \, dx = \mathbb{E}_f[h(X)]$$

▷ Is valid by the Strong Law of Large Numbers

#### ► When $h^2(X)$ has a finite expectation under $f$ ,



$$\frac{\bar{h}_n - \mathbb{E}_f[h(X)]}{\sqrt{v_n}} \rightarrow \mathcal{N}(0, 1)$$

▷ Follows from the Central Limit Theorem

▷  $v_n = \frac{1}{n^2} \sum_{j=1}^n [h(x_j) - \bar{h}_n]^2.$

## Classical Monte Carlo Integration A First Example

► Look at the function




►  $h(x) = [\cos(50x) + \sin(20x)]^2$

► Monitoring Convergence

► R `code`



## Classical Monte Carlo Integration A Caution



- The confidence band produced in this figure is not a 95% confidence band in the classical sense

- They are **Confidence Intervals** were you to stop at a chosen number of iterations

## Classical Monte Carlo Integration Comments

⚡

- ▶ The evaluation of the Monte Carlo error is a bonus
- ▶ It assumes that  $v_n$  is a **proper** estimate of the variance of  $\bar{h}_n$
- ▶ If  $v_n$  does not converge, converges too slowly, a CLT may not apply

## Classical Monte Carlo Integration

### Another Example

#### ► Normal Probability

$$\hat{\Phi}(t) = \frac{1}{n} \sum_{i=1}^n \mathbb{I}_{x_i \leq t} \rightarrow \Phi(t) = \int_{-\infty}^t \frac{1}{\sqrt{2\pi}} e^{-y^2/2} dy$$

▷ The exact variance  $\Phi(t)[1 - \Phi(t)]/n$

▷ Conservative:  $\text{Var} \approx 1/4n$

▷ For a precision of four decimals

▷ Want  $2 \times \sqrt{1/4n} \leq 10^{-4}$  simulations

▷ Take  $n = (10^4)^2 = 10^8$

► This method breaks for tail probabilities

## Importance Sampling Introduction

► **Importance sampling** is based on an alternative formulation of the SLLN

$$\mathbb{E}_f[h(X)] = \int_{\mathcal{X}} h(x) \frac{f(x)}{g(x)} g(x) \, dx = \mathbb{E}_g \left[ \frac{h(X)f(X)}{g(X)} \right] ;$$

- ▷  $f$  is the target density
- ▷  $g$  is the candidate density
  
- ▷ Sound Familiar?

## Importance Sampling Introduction

- Importance sampling is based on an alternative formulation of the SLLN

$$\mathbb{E}_f[h(X)] = \int_{\mathcal{X}} h(x) \frac{f(x)}{g(x)} g(x) \, dx = \mathbb{E}_g \left[ \frac{h(X)f(X)}{g(X)} \right] ;$$

- ▷  $f$  is the target density
- ▷  $g$  is the candidate density
- ▷ Sound Familiar? – Just like Accept–Reject

- So

$$\frac{1}{n} \sum_{j=1}^n \frac{f(X_j)}{g(X_j)} h(X_j) \rightarrow \mathbb{E}_f[h(X)]$$

- As long as

- ▷  $\text{Var}(h(X)f(X)/g(X)) < \infty$
- ▷  $\text{supp}(g) \supset \text{supp}(h \times f)$

## Importance Sampling Revisiting Normal Tail Probabilities

- ▶  $Z \sim \mathcal{N}(0, 1)$  and we are interested in the probability  $P(Z > 4.5)$
- ▶ 

```
> pnorm(-4.5, log=T)
```

```
[1] -12.59242
```
- ▶ Simulating  $Z^{(i)} \sim \mathcal{N}(0, 1)$  only produces a hit once in about 3 million iterations!
  - ▷ Very rare event for the normal
  - ▷ Not-so-rare for a distribution sitting out there!
- ▶ Take  $g = \mathcal{E}xp(1)$  truncated at 4.5:

$$g(y) = \frac{e^{-y}}{\int_{4.5}^{\infty} e^{-x} dx} = e^{-(y-4.5)},$$

- ▶ The IS estimator is

$$\frac{1}{n} \sum_{i=1}^n \frac{f(Y^{(i)})}{g(Y^{(i)})} = \frac{1}{n} \sum_{i=1}^n \frac{e^{-Y_i^2/2 + Y_i - 4.5}}{\sqrt{2\pi}}$$

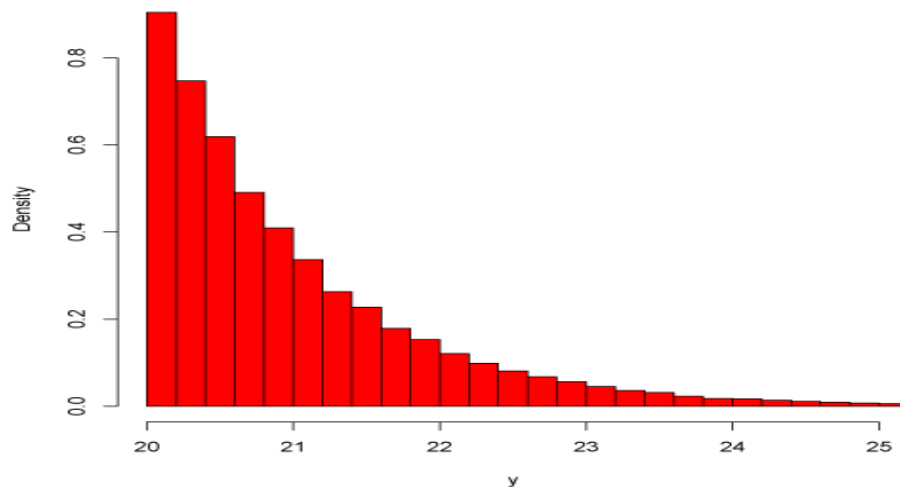
R code

## Importance Sampling Normal Tail Variables

- ▶ The Importance sampler does not give us a sample  $\Rightarrow$  Can use Accept–Reject
- ▶ Sample  $Z \sim \mathcal{N}(0, 1)$ ,  $Z > a \Rightarrow$  Use Exponential Candidate

$$\frac{\frac{1}{\sqrt{2\pi}} \exp(-.5x^2)}{\exp(-(x-a))} = \frac{1}{\sqrt{2\pi}} \exp(-.5x^2 + x + a) \leq \frac{1}{\sqrt{2\pi}} \exp(-.5a^{*2} + a^* + a)$$

▷ Where  $a^* = \max\{a, 1\}$



- ▶ Normals  $> 20$
- ▶ The Twilight Zone
- ▶ R code

## Importance Sampling Comments

- ⚡ Importance sampling has little restriction on the choice of the candidate
  - ▶  $g$  can be chosen from distributions that are easy to simulate
    - ▷ Or efficient in the approximation of the integral.
  - ▶ Moreover, the same sample (generated from  $g$ ) can be used repeatedly
    - ▷ Not only for different functions  $h$  but also for different densities  $f$ .



## Importance Sampling Easy Model - Difficult Distribution

### Example: Beta posterior importance approximation

- Have an observation  $x$  from a beta  $\mathcal{B}(\alpha, \beta)$  distribution,

$$x \sim \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} x^{\alpha-1} (1 - x)^{\beta-1} \mathbb{I}_{[0,1]}(x)$$

- There exists a family of conjugate priors on  $(\alpha, \beta)$  of the form

$$\pi(\alpha, \beta) \propto \left\{ \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \right\}^{\lambda} x_0^{\alpha} y_0^{\beta},$$

where  $\lambda, x_0, y_0$  are hyperparameters,

- The posterior is then equal to

$$\pi(\alpha, \beta | x) \propto \left\{ \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \right\}^{\lambda+1} [xx_0]^{\alpha} [(1 - x)y_0]^{\beta}.$$

## Importance Sampling

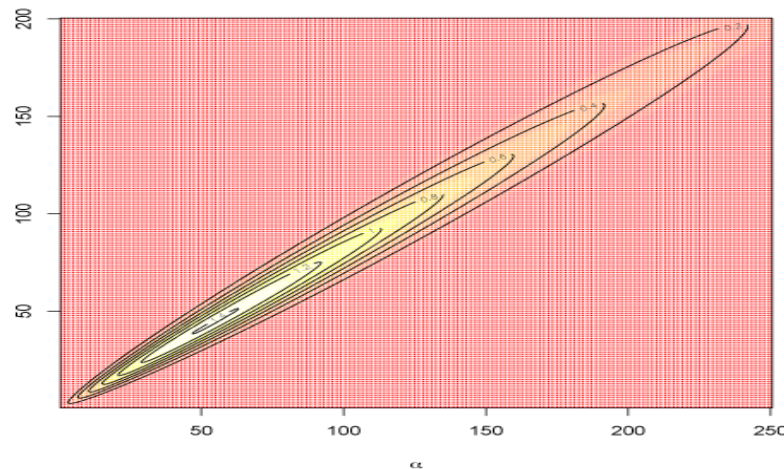
### Easy Model - Difficult Distribution -2

- The posterior distribution is intractable

$$\pi(\alpha, \beta | x) \propto \left\{ \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \right\}^{\lambda+1} [xx_0]^\alpha [(1-x)y_0]^\beta.$$

- ▷ Difficult to deal with the gamma functions
- ▷ Simulating directly from  $\pi(\alpha, \beta | x)$  is impossible.

- What candidate to use?



- Contour Plot
- Suggest a candidate?
- R code

Importance Sampling  
Easy Model - Difficult Distribution – 3

- ▶ Try a Bivariate Student's  $T$  (or Normal)
- ▶ Trial and error
  - ▷ Student's  $\mathcal{T}(3, \mu, \Sigma)$  distribution with  $\mu = (50, 45)$  and

$$\Sigma = \begin{pmatrix} 220 & 190 \\ 190 & 180 \end{pmatrix}$$

- ▷ Produce a reasonable fit
  - ▷ R code
- ▶ Note that we are using the fact that

$$X \sim f(x) \Rightarrow \Sigma^{1/2}X + \mu \sim f\left((x - \mu)' \Sigma^{-1}(x - \mu)\right)$$

## Importance Sampling

### Easy Model - Difficult Distribution – Posterior Means

- The posterior mean of  $\alpha$  is

$$\int \int \alpha \pi(\alpha, \beta | x) d\alpha d\beta = \int \int \left[ \alpha \frac{\pi(\alpha, \beta | x)}{g(\alpha, \beta)} \right] g(\alpha, \beta) d\alpha d\beta \approx \frac{1}{M} \sum_{i=1}^M \alpha_i \frac{\pi(\alpha_i, \beta_i | x)}{g(\alpha_i, \beta_i)}$$

where

$$\triangleright \pi(\alpha, \beta | x) \propto \left\{ \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \right\}^{\lambda+1} [xx_0]^\alpha [(1-x)y_0]^\beta$$

$$\triangleright g(\alpha, \beta) = \mathcal{T}(3, \mu, \Sigma)$$

- Note that  $\pi(\alpha, \beta | x)$  is not normalized, so we have to calculate

$$\frac{\int \int \alpha \pi(\alpha, \beta | x) d\alpha d\beta}{\int \int \pi(\alpha, \beta | x) d\alpha d\beta} \approx \frac{\sum_{i=1}^M \alpha_i \frac{\pi(\alpha_i, \beta_i | x)}{g(\alpha_i, \beta_i)}}{\sum_{i=1}^M \frac{\pi(\alpha_i, \beta_i | x)}{g(\alpha_i, \beta_i)}}$$

- The same samples can be used for every posterior expectation

- R code

## Importance Sampling Probit Analysis

### Example: Probit posterior importance sampling approximation

- $y$  are binary variables, and we have covariates  $x \in \mathbb{R}^p$  such that

$$\Pr(y = 1|x) = 1 - \Pr(y = 0|x) = \Phi(x^T \beta), \quad \beta \in \mathbb{R}^p.$$

- We return to the dataset `Pima.tr`,  $x$ =BMI
- A GLM estimation of the model is (using centered  $x$ )

```
>glm(formula = y ~ x, family = binomial(link = "probit"))
```

Coefficients:

|             | Estimate | Std. Error | z value | Pr(> z ) |     |
|-------------|----------|------------|---------|----------|-----|
| (Intercept) | -0.44957 | 0.09497    | -4.734  | 2.20e-06 | *** |
| x           | 0.06479  | 0.01615    | 4.011   | 6.05e-05 | *** |
| ---         |          |            |         |          |     |

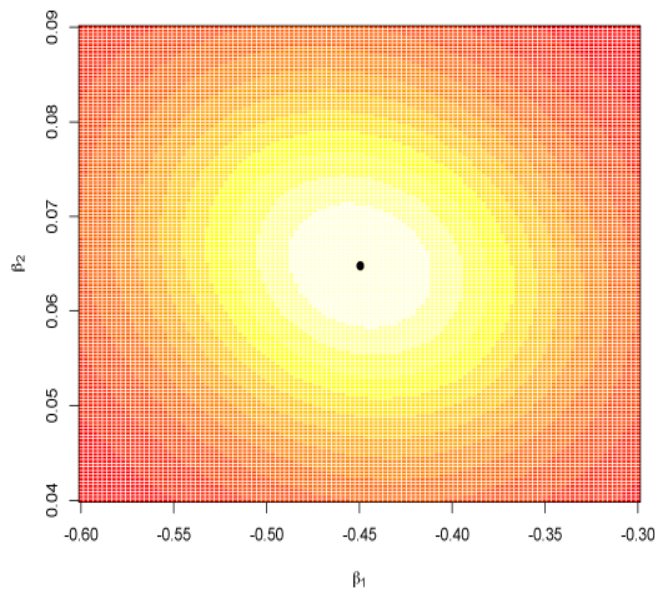
Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

So BMI has a significant impact on the possible presence of diabetes.

## Importance Sampling Bayesian Probit Analysis

- ▶ From a Bayesian perspective, we use a vague prior
  - ▷  $\beta = (\beta_1, \beta_2)$  , each having a  $\mathcal{N}(0, 100)$  distribution
- ▶ With  $\Phi$  the normal cdf, the posterior is proportional to

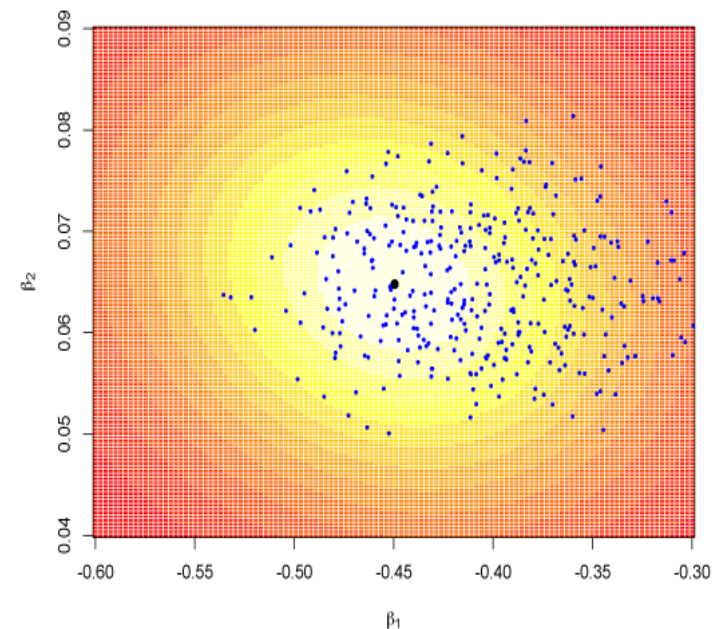
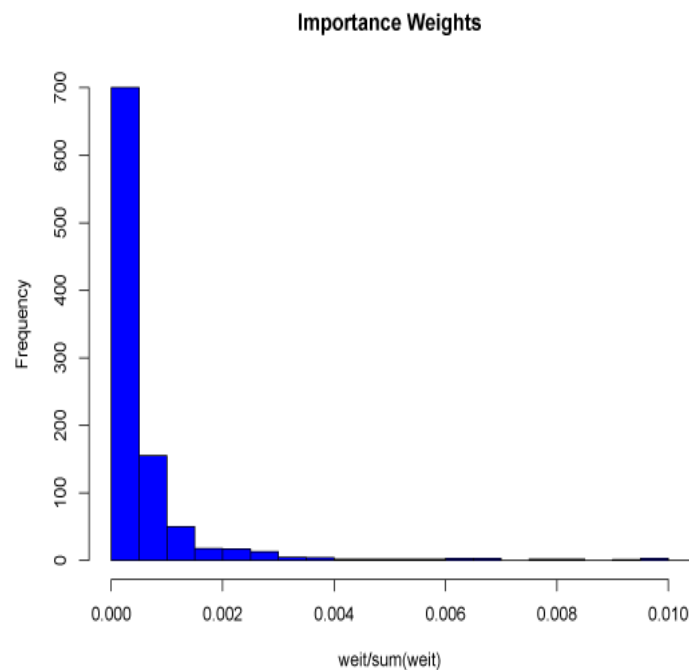
$$\prod_{i=1}^n [\Phi(\beta_1 + (x_i - \bar{x})\beta_2)]^{y_i} [\Phi(-\beta_1 - (x_i - \bar{x})\beta_2)]^{1-y_i} \times e^{-\frac{\beta_1^2 + \beta_2^2}{2 \times 100}}$$



- ▶ Level curves of posterior
- ▶ MLE in the center
- ▶ R [code](#)

## Importance Sampling Probit Analysis Importance Weights

- ▶ Normal candidate centered at the MLE - no finite variance guarantee
- ▶ The importance weights are rather uneven, if not degenerate



- ▶ Right side = reweighted candidate sample (**R** code)
- ▶ Somewhat of a failure

## Chapter 5: Monte Carlo Optimization

*“He invented a game that allowed players to predict the outcome?”*

**Susanna Gregory**

*To Kill or Cure*

### **This Chapter**

- ▶ Two uses of computer-generated random variables to solve optimization problems.
- ▶ **The first use** is to produce stochastic search techniques
  - ▷ To reach the maximum (or minimum) of a function
  - ▷ Avoid being trapped in local maxima (or minima)
  - ▷ Are sufficiently attracted by the global maximum (or minimum).
- ▶ **The second** use of simulation is to approximate the function to be optimized.